

# Biostat 203B Homework 5

Due Mar 20 @ 11:59PM

AUTHOR

Zhiyuan Yu 906405523

## Predicting ICU duration [🔗](#)

Using the ICU cohort `mimiciv_icu_cohort.rds` you built in Homework 4, develop at least three machine learning approaches (logistic regression with enet regularization, random forest, boosting, SVM, MLP, etc) plus a model stacking approach for predicting whether a patient's ICU stay will be longer than 2 days. You should use the `los_long` variable as the outcome. Your algorithms can use patient demographic information (gender, age at ICU `intime`, marital status, race), ICU admission information (first care unit), the last lab measurements before the ICU stay, and first vital measurements during ICU stay as features. You are welcome to use any feature engineering techniques you think are appropriate; but make sure to not use features that are not available at an ICU stay's `intime`. For instance, `last_careunit` cannot be used in your algorithms.

```
library(gtsummary)
library(tidyverse)
```

— Attaching core tidyverse packages — tidyverse 2.0.0 —

```
✓ dplyr      1.1.4    ✓ readr      2.1.5
✓ forcats    1.0.0    ✓ stringr    1.5.1
✓ ggplot2    3.5.1    ✓ tibble     3.2.1
✓ lubridate  1.9.4    ✓ tidyr      1.3.1
✓ purrr      1.0.4
```

— Conflicts — tidyverse\_conflicts() —

```
✗ dplyr::filter() masks stats::filter()
```

```
✗ dplyr::lag()     masks stats::lag()
```

ℹ Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

```
library(tidymodels)
```

— Attaching packages — tidymodels 1.3.0 —

```
✓ broom       1.0.7    ✓ rsample     1.2.1
✓ dials       1.4.0    ✓ tune        1.3.0
✓ infer       1.0.7    ✓ workflows   1.2.0
✓ modeldata   1.4.0    ✓ workflowsets 1.1.0
✓ parsnip     1.3.0    ✓ yardstick   1.3.2
✓ recipes     1.1.1
```

— Conflicts — tidymodels\_conflicts() —

```
✗ scales::discard() masks purrr::discard()
```

```
✗ dplyr::filter()   masks stats::filter()
```

```
✗ recipes::fixed()  masks stringr::fixed()
```

```
✗ dplyr::lag()      masks stats::lag()
```

✗ yardstick::spec() masks readr::spec()  
✗ recipes::step() masks stats::step()

```
library(dplyr)
library(haven)
library(recipes)
library(GGally)
```

Registered S3 method overwritten by 'GGally':  
method from  
+.gg ggplot2

```
library(ranger)
library(xgboost)
```

Attaching package: 'xgboost'

The following object is masked from 'package:dplyr':

slice

```
library(stacks)
library(yardstick)
library(purrr)
library(vip)
```

Attaching package: 'vip'

The following object is masked from 'package:utils':

vi

```
library(parsnip)
library(tune)
library(dials)
library(purrr)
```

## 1. Data preprocessing and feature engineering.

**Solution:** I put the mimic\_icu\_cohort.rds file from HW4 to my current working directory of HW5. Step 1: Check for missing values

```
mimic_icu_cohort <- read_rds("mimic_icu_cohort.rds") %>%
  filter(!is.na(los_long)) %>%
  select(-last_careunit, -dod, -discharge_location, -hospital_expire_flag, -los,
         -intime, -outtime_x, -admittime, -disctime,
```

```

      -deathtime, -edregtime, -edouttime, -outtime_y, -admit_provider_id) %>%
mutate(
  los_long = as.factor(los_long),
  insurance = as.factor(insurance),
  marital_status = as.factor(marital_status),
  language = as.factor(language)
)

# Check for missing values
colSums(is.na(mimic_icu_cohort))

```

subject_id	0	hadm_id	0
stay_id	0	first_careunit	0
admission_type	0	admission_location	0
insurance	1523	language	396
marital_status	7756	race	0
gender	0	anchor_age	0
anchor_year	0	anchor_year_group	0
bicarbonate	11549	chloride	11351
creatinine	8027	glucose	11654
potassium	11387	sodium	11330
hematocrit	6751	white_blood_cells	6850
heart_rate	86	systolic_non_invasive_blood_pressure	1370
diastolic_non_invasive_blood_pressure	1375	temperature_fahrenheit	1675
respiratory_rate	198	age_intime	0
los_long	0		

Step 2: Take a peek at the data types

```
sapply(mimic_icu_cohort, class)
```

subject_id	"integer"	hadm_id	"integer"
stay_id	"integer"	first_careunit	"factor"
admission_type		admission_location	

"factor"	"factor"
insurance	language
"factor"	"factor"
marital_status	race
"factor"	"factor"
gender	anchor_age
"character"	"integer"
anchor_year	anchor_year_group
"integer"	"character"
bicarbonate	chloride
"numeric"	"numeric"
creatinine	glucose
"numeric"	"numeric"
potassium	sodium
"numeric"	"numeric"
hematocrit	white_blood_cells
"numeric"	"numeric"
heart_rate	systolic_non_invasive_blood_pressure
"numeric"	"numeric"
diastolic_non_invasive_blood_pressure	temperature_fahrenheit
"numeric"	"numeric"
respiratory_rate	age_intime
"numeric"	"integer"
los_long	
"factor"	

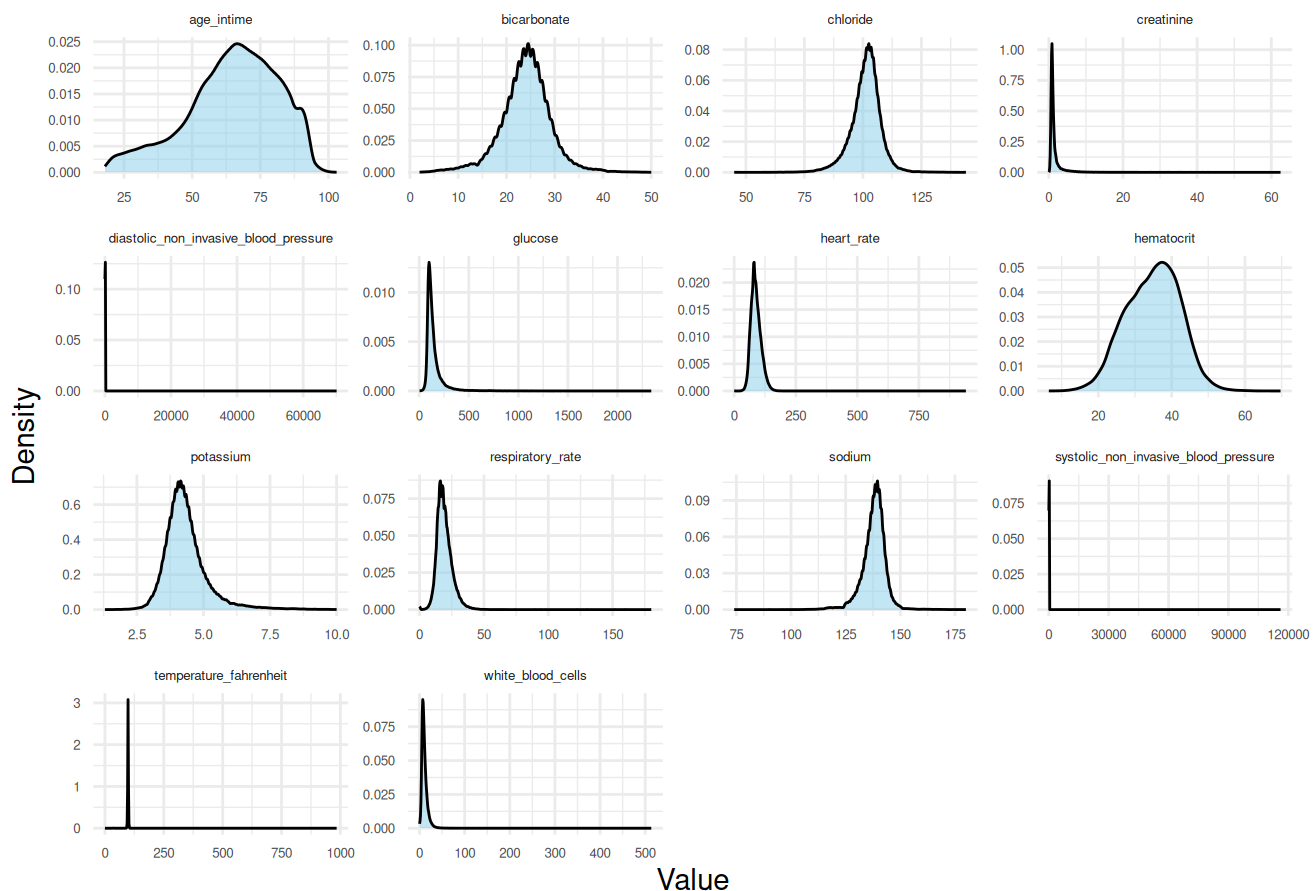
Step 3: Create histograms for the continuous variables to see their distributions

```
# Define a vector with the names of the continuous variables
cont_vars <- c(
  "bicarbonate", "chloride", "creatinine", "glucose", "potassium", "sodium",
  "hematocrit", "white_blood_cells", "heart_rate",
  "systolic_non_invasive_blood_pressure",
  "diastolic_non_invasive_blood_pressure", "temperature_fahrenheit",
  "respiratory_rate", "age_intime"
)

# Reshape the dataset into long format
hist_data_alt <- mimic_icu_cohort %>%
  select(all_of(cont_vars)) %>%
  pivot_longer(
    cols = everything(),
    names_to = "variable",
    values_to = "value"
  )

# Plot the density distributions for each variable
ggplot(hist_data_alt, aes(x = value)) +
  geom_density(kernel = "cosine", fill = "skyblue", alpha = 0.5, na.rm = TRUE) +
  facet_wrap(~ variable, scales = "free") +
  labs(x = "Value", y = "Density") +
```

```
theme_minimal() +
theme(
  plot.margin = margin(12, 12, 12, 12, "pt"),
  strip.text = element_text(size = 5),
  axis.text.x = element_text(size = 5),
  axis.text.y = element_text(size = 5)
)
```



It is clear that most of the data distributions are either left or right skewed in different levels, therefore I will use median instead of mean to impute the data.

**2. Partition data into 50% training set and 50% test set. Stratify partitioning according to `los_long`. For grading purpose, sort the data by `subject_id`, `hadm_id`, and `stay_id` and use the seed 203 for the initial data split. Below is the sample code.** [🔗](#)

**Solution:** Step 1: split the data into training set and test set

```
set.seed(203)

# sort the data by subject_id, hadm_id, and stay_id
mimic_icu_cohort <- mimic_icu_cohort |>
  arrange(subject_id, hadm_id, stay_id) |>
```

```
select(-subject_id, -hadm_id, -stay_id)

# partition data into 50% training set and 50% test set
data_split <- initial_split(
  mimim_icu_cohort,
  strata = "los_long",
  prop = 0.5
)

data_split
```

<Training/Testing/Total>  
<47221/47223/94444>

```
# check the training set
train_data <- training(data_split)
dim(train_data)
```

[1] 47221 26

```
# check the training and testing set
test_data <- testing(data_split)
dim(test_data)
```

[1] 47223 26

Step 2: Preprocess the data

```
recipe <- recipe(los_long ~ ., data = train_data) %>%
  step_impute_median(all_of(c("bicarbonate", "chloride", "creatinine",
    "glucose", "potassium",
    "sodium", "hematocrit", "white_blood_cells",
    "heart_rate",
    "systolic_non_invasive_blood_pressure",
    "diastolic_non_invasive_blood_pressure",
    "temperature_fahrenheit", "respiratory_rate",
    "age_intime")))) %>%
  step_impute_mode(insurance, marital_status, language) %>%
  step_dummy(all_nominal_predictors()) %>%
  step_zv(all_numeric_predictors()) %>%
  step_normalize(all_numeric_predictors()) %>%
  print()
```

— Recipe —

— Inputs

Number of variables by role

outcome: 1  
 predictor: 25

#### — Operations

- Median imputation for: `all_of(c("bicarbonate", "chloride", "creatinine", "glucose", "potassium", "sodium", "hematocrit", "white_blood_cells", "heart_rate", "systolic_non_invasive_blood_pressure", "diastolic_non_invasive_blood_pressure", "temperature_fahrenheit", "respiratory_rate", "age_intime"))`
- Mode imputation for: `insurance, marital_status, language`
- Dummy variables from: `all_nominal_predictors()`
- Zero variance filter on: `all_numeric_predictors()`
- Centering and scaling for: `all_numeric_predictors()`

### 3. Train and tune the models using the training set.

First approach: Logistic regression with elastic net regularization

Step 1: Define recipe

```
logit_recipe <- recipe
```

Step 1: Define the logistic regression model for classification and set the engine

```
logit_model <- logistic_reg(
  penalty = tune(),
  mixture = tune()
) %>%
  set_engine("glmnet", standardize = FALSE)

print(logit_model)
```

Logistic Regression Model Specification (classification)

Main Arguments:

```
penalty = tune()
mixture = tune()
```

Engine-Specific Arguments:

```
standardize = FALSE
```

Computational engine: `glmnet`

## Step 2: Define workflow

```
logit_workflow <- workflow() %>%
  add_recipe(logit_recipe) %>%
  add_model(logit_model) %>%
  print()
```

## == Workflow ==

Preprocessor: Recipe

Model: logistic\_reg()

## — Preprocessor —

5 Recipe Steps

- step\_impute\_median()
- step\_impute\_mode()
- step\_dummy()
- step\_zv()
- step\_normalize()

## — Model —

Logistic Regression Model Specification (classification)

Main Arguments:

```
penalty = tune()
mixture = tune()
```

Engine-Specific Arguments:

```
standardize = FALSE
```

Computational engine: glmnet

## Step 3: Tune the grid and do cross-validation folds

```
logit_param_grid <- grid_regular(
  penalty(range = c(-4, 1)),
  mixture(),
  levels = c(100, 5)) %>%
  print()
```

# A tibble: 500 × 2

	penalty	mixture
	<dbl>	<dbl>
1	0.0001	0
2	0.000112	0
3	0.000126	0
4	0.000142	0
5	0.000159	0
6	0.000179	0



```

7 0.000201      0
8 0.000226      0
9 0.000254      0
10 0.000285     0
# i 490 more rows

```

```

set.seed(203)

# define the number of folds for cross-validation is 5
logit_folds <- vfold_cv(train_data, v = 5)
logit_folds

```

```

# 5-fold cross-validation
# A tibble: 5 × 2
  splits      id
  <list>     <chr>
1 <split [37776/9445]> Fold1
2 <split [37777/9444]> Fold2
3 <split [37777/9444]> Fold3
4 <split [37777/9444]> Fold4
5 <split [37777/9444]> Fold5

```

Step 4: Fit the corss-validated models and select the best model

```

suppressMessages(suppressWarnings({
  if (file.exists("logit_fit.rds")) {
    logit_fit <- read_rds("logit_fit.rds")
    logit_fit

  } else {
    (logit_fit <- logit_workflow |>
      tune_grid(
        resamples = logit_folds,
        grid = logit_param_grid,
        metrics = metric_set(roc_auc, accuracy)
      ) |>
      system.time()

    logit_fit |>
      write_rds("logit_fit.rds")

    logit_fit
  }
}))

```

```

# Tuning results
# 5-fold cross-validation
# A tibble: 5 × 5
  splits      id .metrics      .notes      .predictions
  <list>     <chr> <list>      <list>      <list>

```

```

1 <split [37776/9445]> Fold1 <tibble [1,000 × 6]> <tibble [0 × 3]> <tibble>
2 <split [37777/9444]> Fold2 <tibble [1,000 × 6]> <tibble [0 × 3]> <tibble>
3 <split [37777/9444]> Fold3 <tibble [1,000 × 6]> <tibble [0 × 3]> <tibble>
4 <split [37777/9444]> Fold4 <tibble [1,000 × 6]> <tibble [0 × 3]> <tibble>
5 <split [37777/9444]> Fold5 <tibble [1,000 × 6]> <tibble [0 × 3]> <tibble>

```

Step 5: Visualize CV result

```

logit_fit |>
  # aggregate metrics from 5-fold cross-validation
  collect_metrics() |>
  print(width = Inf) |>
  filter(.metric == "roc_auc") |>
  ggplot(mapping = aes(x = penalty, y = mean,
                       color = factor(mixture))) +
  geom_point() +
  labs(x = "Penalty", y = "CV AUC") +
  scale_x_log10()

```


# A tibble: 1,000 × 8

	penalty	mixture	.metric	.estimator	mean	n	std_err
	<dbl>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>
1	0.0001	0	accuracy	binary	0.579	5	0.00200
2	0.0001	0	roc_auc	binary	0.610	5	0.00149
3	0.000112	0	accuracy	binary	0.579	5	0.00200
4	0.000112	0	roc_auc	binary	0.610	5	0.00149
5	0.000126	0	accuracy	binary	0.579	5	0.00200
6	0.000126	0	roc_auc	binary	0.610	5	0.00149
7	0.000142	0	accuracy	binary	0.579	5	0.00200
8	0.000142	0	roc_auc	binary	0.610	5	0.00149
9	0.000159	0	accuracy	binary	0.579	5	0.00200
10	0.000159	0	roc_auc	binary	0.610	5	0.00149

```

.config
<chr>
1 Preprocessor1_Model001
2 Preprocessor1_Model001
3 Preprocessor1_Model002
4 Preprocessor1_Model002
5 Preprocessor1_Model003
6 Preprocessor1_Model003
7 Preprocessor1_Model004
8 Preprocessor1_Model004
9 Preprocessor1_Model005
10 Preprocessor1_Model005
# i 990 more rows

```



Step 6: Show the top 5 models

```
logit_fit %>%
  show_best(metric = "roc_auc")
```

# A tibble: 5 × 8

	penalty	mixture	.metric	.estimator	mean	n	std_err	.config
	<dbl>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	0.000911	1	roc_auc	binary	0.610	5	0.00137	Preprocessor1_Model420
2	0.000811	1	roc_auc	binary	0.610	5	0.00139	Preprocessor1_Model419
3	0.00102	1	roc_auc	binary	0.610	5	0.00135	Preprocessor1_Model421
4	0.000722	1	roc_auc	binary	0.610	5	0.00141	Preprocessor1_Model418
5	0.000643	1	roc_auc	binary	0.610	5	0.00143	Preprocessor1_Model417

Step 7: Select the best model

```
best_logit <- logit_fit |>
  select_best(metric = "roc_auc")
best_logit
```

# A tibble: 1 × 3

penalty	mixture	.config
---------	---------	---------

```
<dbl> <dbl> <chr>
```

```
1 0.000911      1 Preprocessor1_Model1420
```

Step 8: Finalize model

```
# Final workflow
logit_final_workflow <- logit_workflow %>%
  finalize_workflow(best_logit)
logit_final_workflow
```

== Workflow ==

Preprocessor: Recipe

Model: logistic\_reg()

— Preprocessor —

5 Recipe Steps

- step\_impute\_median()
- step\_impute\_mode()
- step\_dummy()
- step\_zv()
- step\_normalize()

— Model —

Logistic Regression Model Specification (classification)

Main Arguments:

```
penalty = 0.000911162756115489
mixture = 1
```

Engine-Specific Arguments:

```
standardize = FALSE
```

Computational engine: glmnet

```
# Fit the whole training set, then predict the test cases
logit_final_fit <- logit_final_workflow %>%
  last_fit(data_split)
```

New names:

New names:

- `anchor\_year\_group\_X2011...2013` -> `anchor\_year\_group\_X2011`
- `anchor\_year\_group\_X2014...2016` -> `anchor\_year\_group\_X2014`
- `anchor\_year\_group\_X2017...2019` -> `anchor\_year\_group\_X2017`
- `anchor\_year\_group\_X2020...2022` -> `anchor\_year\_group\_X2020`

```
logit_final_fit
```

# Resampling results

# Manual resampling

```
# A tibble: 1 × 6
```

	splits	id	.metrics	.notes	.predictions	.workflow
	<list>	<chr>	<list>	<list>	<list>	<list>
1	<split [47221/47223]>	train/test sp...	<tibble>	<tibble>	<tibble>	<workflow>

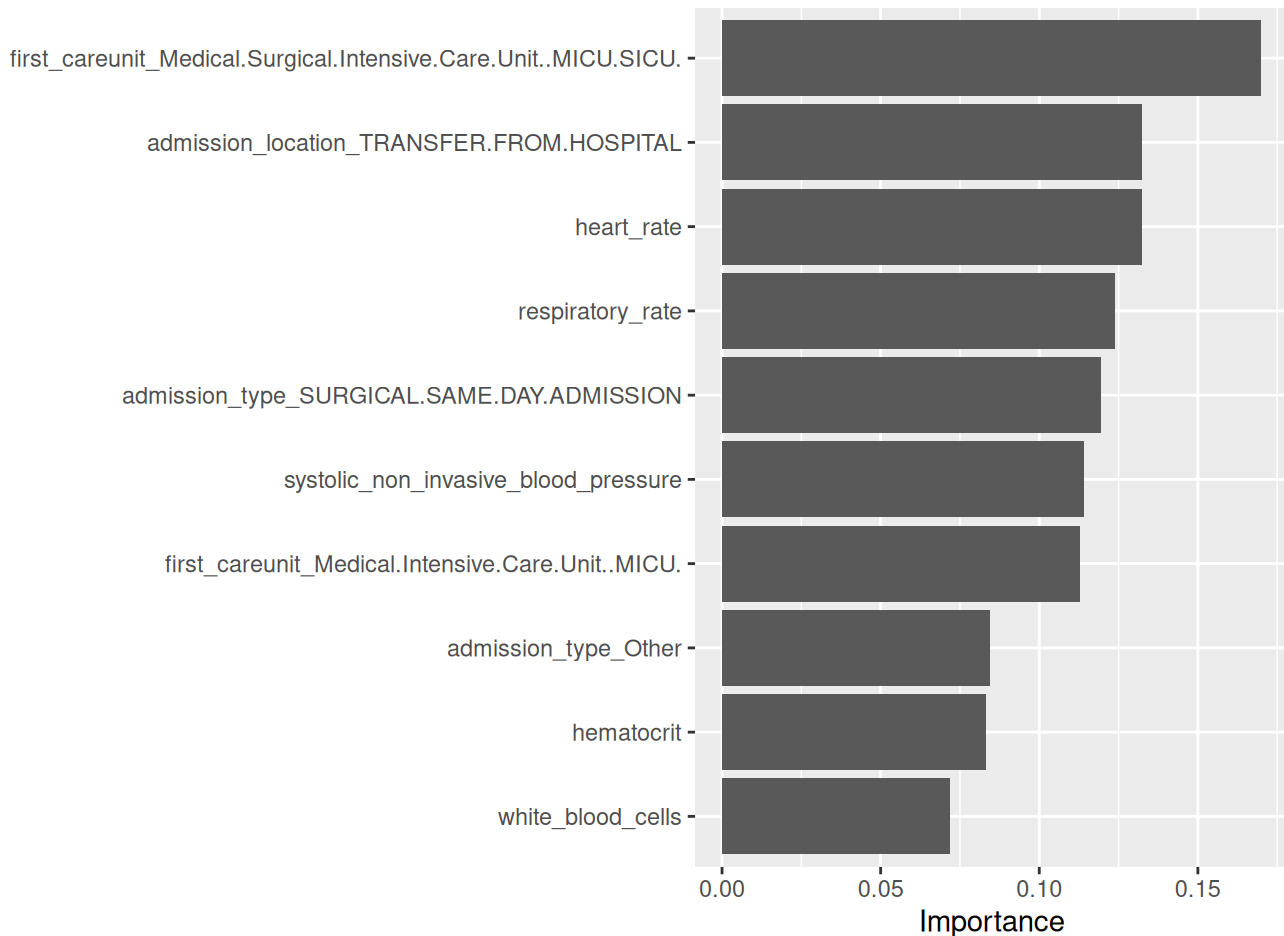
```
# Test metrics
logit_final_fit %>%
  collect_metrics()
```

```
# A tibble: 3 × 4
```

	.metric	.estimator	.estimate	.config
	<chr>	<chr>	<dbl>	<chr>
1	accuracy	binary	0.578	Preprocessor1_Model1
2	roc_auc	binary	0.608	Preprocessor1_Model1
3	brier_class	binary	0.241	Preprocessor1_Model1

Step 9: Plot the variable importance

```
logit_final_fit %>%
  extract_fit_parsnip() %>%
  vip::vip() %>%
  print()
```



Summary of the logistic regression with elastic net regularization: Based on the final logistic regression model, its accuracy is 0.5782775, meaning that the model only has 57.8% rate of correct prediction on the test set. The model also has AUC of 0.60881425, indicating that the model can classify 60.9% of the test set correctly.

The 10 most important predictors are shown in the variable importance plot, and the most important predictor is the first care unit.

## Second Approach: Random Forest

### Step 1: Define recipe

```
# define the recipe
rf_recipe <- recipe
```

### Step 2: Define Random Forest Model

```
rf_model <- rand_forest(
  mode = "classification",
  mtry = tune(),
  trees = tune()
) %>%
  set_engine("ranger", importance = "impurity")
rf_model
```

#### Random Forest Model Specification (classification)

##### Main Arguments:

```
mtry = tune()
trees = tune()
```

##### Engine-Specific Arguments:

```
importance = impurity
```

Computational engine: ranger

### Step 3: Define workflow

```
rf_workflow <- workflow() %>%
  add_recipe(rf_recipe) %>%
  add_model(rf_model) %>%
  print()
```

== Workflow ==

Preprocessor: Recipe

Model: rand\_forest()

— Preprocessor —

5 Recipe Steps

- `step_impute_median()`
- `step_impute_mode()`
- `step_dummy()`
- `step_zv()`
- `step_normalize()`

— Model —

---

### Random Forest Model Specification (classification)

#### Main Arguments:

```
mtry = tune()
trees = tune()
```

#### Engine-Specific Arguments:

```
importance = impurity
```

Computational engine: ranger

#### Step 4: Tuning grid

```
rf_param_grid <- grid_regular(
  trees(range = c(100L, 500L)),
  mtry(range = c(2L, 5L)),
  levels = c(5, 5)
)
```

```
rf_param_grid
```

# A tibble: 20 × 2

	trees	mtry
	<int>	<int>
1	100	2
2	200	2
3	300	2
4	400	2
5	500	2
6	100	3
7	200	3
8	300	3
9	400	3
10	500	3
11	100	4
12	200	4
13	300	4
14	400	4
15	500	4
16	100	5
17	200	5
18	300	5
19	400	5
20	500	5

## Step 5: Cross-validation

```
# Set cross-validation partitions
set.seed(203)

rf_folds <- vfold_cv(train_data, v = 5)
rf_folds
```

```
# 5-fold cross-validation
# A tibble: 5 × 2
  splits      id
  <list>     <chr>
1 <split [37776/9445]> Fold1
2 <split [37777/9444]> Fold2
3 <split [37777/9444]> Fold3
4 <split [37777/9444]> Fold4
5 <split [37777/9444]> Fold5
```

```
# Fit cross-validation
if (file.exists("rf_fit.rds")) {
  rf_fit <- read_rds("rf_fit.rds")
  rf_fit

} else {
  (rf_fit <- rf_workflow %>%
    tune_grid(
      resamples = rf_folds,
      grid = rf_param_grid,
      metrics = metric_set(roc_auc, accuracy),
      control = control_stack_grid()
    )) %>%
  system.time()

  rf_fit %>%
    write_rds("rf_fit.rds")

  rf_fit
}
```

```
# Tuning results
# 5-fold cross-validation
# A tibble: 5 × 5
  splits      id .metrics      .notes      .predictions
  <list>     <chr> <list>      <list>      <list>
1 <split [37776/9445]> Fold1 <tibble [40 × 6]> <tibble [0 × 3]> <tibble>
2 <split [37777/9444]> Fold2 <tibble [40 × 6]> <tibble [0 × 3]> <tibble>
3 <split [37777/9444]> Fold3 <tibble [40 × 6]> <tibble [0 × 3]> <tibble>
4 <split [37777/9444]> Fold4 <tibble [40 × 6]> <tibble [0 × 3]> <tibble>
5 <split [37777/9444]> Fold5 <tibble [40 × 6]> <tibble [0 × 3]> <tibble>
```




## Step 6: Visualize CV results

```
rf_fit |>
  collect_metrics() %>%
  print(width = Inf) %>%
  filter(.metric == "roc_auc") %>%
  ggplot(mapping = aes(x = trees, y = mean,
                      color = factor(mtry))) +
  geom_point() +
  labs(x = "Num. of Trees", y = "CV AUC")
```

# A tibble: 40 × 8

	mtry	trees	.metric	.estimator	mean	n	std_err	.config
	<int>	<int>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	2	100	accuracy	binary	0.597	5	0.00253	Preprocessor1_Model01
2	2	100	roc_auc	binary	0.637	5	0.00280	Preprocessor1_Model01
3	2	200	accuracy	binary	0.602	5	0.00249	Preprocessor1_Model02
4	2	200	roc_auc	binary	0.642	5	0.00235	Preprocessor1_Model02
5	2	300	accuracy	binary	0.599	5	0.00218	Preprocessor1_Model03
6	2	300	roc_auc	binary	0.641	5	0.00195	Preprocessor1_Model03
7	2	400	accuracy	binary	0.600	5	0.00196	Preprocessor1_Model04
8	2	400	roc_auc	binary	0.642	5	0.00220	Preprocessor1_Model04
9	2	500	accuracy	binary	0.601	5	0.00237	Preprocessor1_Model05
10	2	500	roc_auc	binary	0.643	5	0.00220	Preprocessor1_Model05

# i 30 more rows



Step 7: Show the top 5 models

```
rf_fit |>
  show_best(metric = "roc_auc")
```

# A tibble: 5 × 8

	mtry	trees	.metric	.estimator	mean	n	std_err	.config
	<int>	<int>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	5	500	roc_auc	binary	0.650	5	0.00194	Preprocessor1_Model120
2	4	300	roc_auc	binary	0.649	5	0.00158	Preprocessor1_Model113
3	4	500	roc_auc	binary	0.649	5	0.00176	Preprocessor1_Model115
4	4	400	roc_auc	binary	0.649	5	0.00194	Preprocessor1_Model114
5	5	400	roc_auc	binary	0.649	5	0.00177	Preprocessor1_Model119

Step 8: Select the best model

```
best_rf <- rf_fit |>
  select_best(metric = "roc_auc")
best_rf
```

# A tibble: 1 × 3

mtry	trees	.config
5	500	Preprocessor1_Model120

```
<int> <int> <chr>
```

```
1      5      500 Preprocessor1_Model120
```

Step 9: Finalize the model

```
# Final workflow
rf_final_workflow <- rf_workflow %>%
  finalize_workflow(best_rf)
rf_final_workflow
```

== Workflow ==

Preprocessor: Recipe

Model: rand\_forest()

— Preprocessor —

5 Recipe Steps

- step\_impute\_median()
- step\_impute\_mode()
- step\_dummy()
- step\_zv()
- step\_normalize()

— Model —

Random Forest Model Specification (classification)

Main Arguments:

```
mtry = 5
```

```
trees = 500
```

Engine-Specific Arguments:

```
importance = impurity
```

Computational engine: ranger

```
# Fit the whole training set, then predict the test cases
if (file.exists("rf_final_fit.rds")) {
  rf_final_fit <- read_rds("rf_final_fit.rds")
  rf_final_fit

} else {

  # fit the final model on the whole training set
  rf_final_fit <- rf_final_workflow %>%
    last_fit(data_split)

  rf_final_fit %>%
    write_rds("rf_final_fit.rds")
  rf_final_fit
}
```


```
# Resampling results
# Manual resampling
# A tibble: 1 × 6
  splits          id      .metrics .notes  .predictions .workflow
  <list>         <chr>    <list>  <list>  <list>       <list>
1 <split [47221/47223]> train/test sp... <tibble> <tibble> <tibble>    <workflow>
```

```
# Test metrics
rf_final_fit %>%
  collect_metrics()
```

```
# A tibble: 3 × 4
  .metric      .estimator .estimate .config
  <chr>        <chr>      <dbl> <chr>
1 accuracy    binary      0.606 Preprocessor1_Model11
2 roc_auc     binary      0.647 Preprocessor1_Model11
3 brier_class binary      0.234 Preprocessor1_Model11
```

Step 10: Plot the variable importance

```
rf_final_fit %>%
  extract_fit_engine() %>%
  vip::vip() %>%
  print()
```



Summary of the random forest model: The model has accuracy of 0.6060818, meaning that the model can predict 60.6% of the test dataset correctly. The AUC of the model is 0.6474515, which means that the final random forest model can correctly classify 64.7% of the test dataset. At this point, we can see that the random forest model has better performance in both accuracy and AUC than the logistic model.

The 10 most important predictors are listed in the variable importance plot. The most important variable is systolic noninvasive blood pressure.

## Third approach: Boosting (XGBoost)

Step 1: Define recipe

```
xgb_recipe <- recipe
```

Step 2: Define boosting model

```
xgb_model <- boost_tree(
  mode = "classification",
  trees = 1000,
  tree_depth = tune(),
  learn_rate = tune()
) %>%
  set_engine("xgboost")

xgb_model
```

Boosted Tree Model Specification (classification)

Main Arguments:

```
trees = 1000
tree_depth = tune()
learn_rate = tune()
```

Computational engine: xgboost

Step 3: Define workflow

```
xgb_workflow <- workflow() %>%
  add_recipe(xgb_recipe) %>%
  add_model(xgb_model)
xgb_workflow
```

== Workflow ==

Preprocessor: Recipe

Model: boost\_tree()

— Preprocessor —

5 Recipe Steps

- `step_impute_median()`
- `step_impute_mode()`
- `step_dummy()`
- `step_zv()`
- `step_normalize()`

— Model —

---

#### Boosted Tree Model Specification (classification)

Main Arguments:

```
trees = 1000
tree_depth = tune()
learn_rate = tune()
```

Computational engine: `xgboost`

Step 4: Tuning grid

```
xgb_param_grid <- grid_regular(
  tree_depth(range = c(1L, 3L)),
  learn_rate(range = c(-3, 0), trans = log10_trans()),
  levels = c(3, 3)
)
xgb_param_grid
```

# A tibble: 9 × 2

	tree_depth	learn_rate
	<int>	<dbl>
1	1	0.001
2	2	0.001
3	3	0.001
4	1	0.0316
5	2	0.0316
6	3	0.0316
7	1	1
8	2	1
9	3	1

Step 5: Cross-validation

```
# Set cross-validation partitions
set.seed(203)

xgb_folds <- vfold_cv(train_data, v = 5)
xgb_folds
```

# 5-fold cross-validation

# A tibble: 5 × 2

splits	id
<list>	<chr>

```

1 <split [37776/9445]> Fold1
2 <split [37777/9444]> Fold2
3 <split [37777/9444]> Fold3
4 <split [37777/9444]> Fold4
5 <split [37777/9444]> Fold5

```

```

# Fit cross-validation
if (file.exists("xgb_fit.rds")) {
  xgb_fit <- read_rds("xgb_fit.rds")
  xgb_fit
} else {
  (xgb_fit <- xgb_workflow %>%
    tune_grid(
      resamples = xgb_folds,
      grid = xgb_param_grid,
      metrics = metric_set(roc_auc, accuracy),
      control = control_stack_grid()
    )) %>%
    system.time()

  xgb_fit %>%
    write_rds("xgb_fit.rds")

  xgb_fit
}

```

# Tuning results

# 5-fold cross-validation

# A tibble: 5 × 5

	splits	id	.metrics	.notes	.predictions
	<list>	<chr>	<list>	<list>	<list>
1	<split [37776/9445]>	Fold1	<tibble [18 × 6]>	<tibble [0 × 3]>	<tibble>
2	<split [37777/9444]>	Fold2	<tibble [18 × 6]>	<tibble [0 × 3]>	<tibble>
3	<split [37777/9444]>	Fold3	<tibble [18 × 6]>	<tibble [0 × 3]>	<tibble>
4	<split [37777/9444]>	Fold4	<tibble [18 × 6]>	<tibble [0 × 3]>	<tibble>
5	<split [37777/9444]>	Fold5	<tibble [18 × 6]>	<tibble [0 × 3]>	<tibble>

Step 6: Visualize CV results

```

xgb_fit |>
  collect_metrics() |>
  print(width = Inf) |>
  filter(.metric == "roc_auc") |>
  ggplot(mapping = aes(x = learn_rate, y = mean,
                       color = factor(tree_depth))) +
  geom_point() +
  labs(x = "Learning Rate", y = "CV AUC") +
  scale_x_log10()

```

# A tibble: 18 × 8

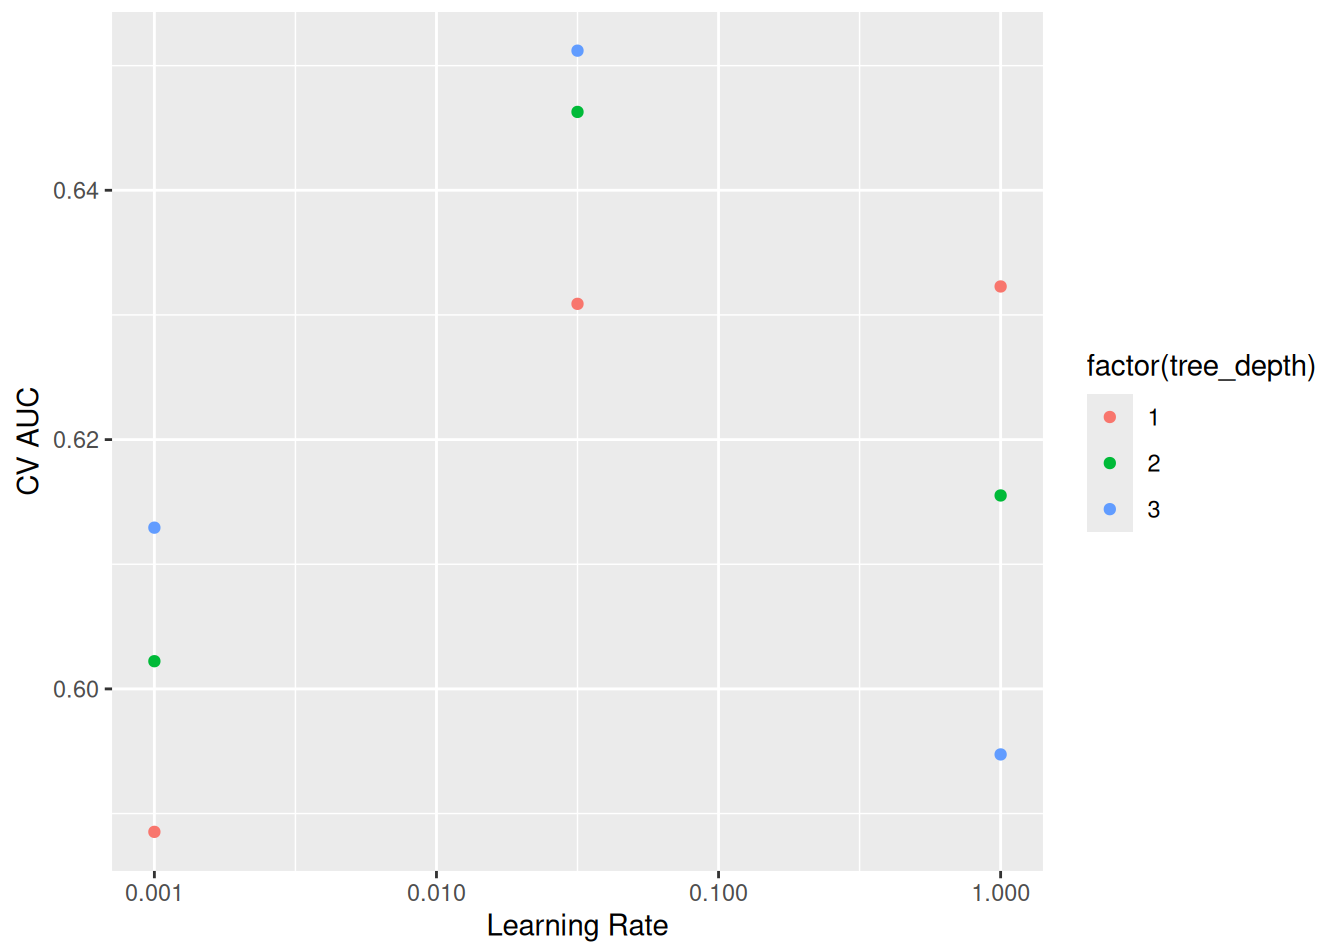
	tree_depth	learn_rate	.metric	.estimator	mean	n	std_err
	<int>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>
1	1	0.001	accuracy	binary	0.559	5	0.00245
2	1	0.001	roc_auc	binary	0.589	5	0.00293
3	2	0.001	accuracy	binary	0.572	5	0.00289
4	2	0.001	roc_auc	binary	0.602	5	0.00302
5	3	0.001	accuracy	binary	0.580	5	0.00236
6	3	0.001	roc_auc	binary	0.613	5	0.00299
7	1	0.0316	accuracy	binary	0.592	5	0.00241
8	1	0.0316	roc_auc	binary	0.631	5	0.00130
9	2	0.0316	accuracy	binary	0.605	5	0.000971
10	2	0.0316	roc_auc	binary	0.646	5	0.00115
11	3	0.0316	accuracy	binary	0.608	5	0.000896
12	3	0.0316	roc_auc	binary	0.651	5	0.00139
13	1	1	accuracy	binary	0.596	5	0.00197
14	1	1	roc_auc	binary	0.632	5	0.000888
15	2	1	accuracy	binary	0.585	5	0.00166
16	2	1	roc_auc	binary	0.616	5	0.00202
17	3	1	accuracy	binary	0.571	5	0.00148
18	3	1	roc_auc	binary	0.595	5	0.00129

.config

<chr>

1 Preprocessor1\_Model11  
 2 Preprocessor1\_Model11  
 3 Preprocessor1\_Model12  
 4 Preprocessor1\_Model12  
 5 Preprocessor1\_Model13  
 6 Preprocessor1\_Model13  
 7 Preprocessor1\_Model14  
 8 Preprocessor1\_Model14  
 9 Preprocessor1\_Model15  
 10 Preprocessor1\_Model15  
 11 Preprocessor1\_Model16  
 12 Preprocessor1\_Model16  
 13 Preprocessor1\_Model17  
 14 Preprocessor1\_Model17  
 15 Preprocessor1\_Model18  
 16 Preprocessor1\_Model18  
 17 Preprocessor1\_Model19  
 18 Preprocessor1\_Model19





Step 7: Show the best 5 models

```
xgb_fit |>
  show_best(metric = "roc_auc")
```

# A tibble: 5 × 8

	tree_depth	learn_rate	.metric	.estimator	mean	n	std_err	.config
	<int>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	3	0.0316	roc_auc	binary	0.651	5	0.00139	Preprocessor1_M...
2	2	0.0316	roc_auc	binary	0.646	5	0.00115	Preprocessor1_M...
3	1	1	roc_auc	binary	0.632	5	0.000888	Preprocessor1_M...
4	1	0.0316	roc_auc	binary	0.631	5	0.00130	Preprocessor1_M...
5	2	1	roc_auc	binary	0.616	5	0.00202	Preprocessor1_M...

Step 8: Select the best model

```
best_xgb <- xgb_fit |>
  select_best(metric = "roc_auc")
best_xgb
```

# A tibble: 1 × 3

tree_depth	learn_rate	.config
3	0.0316	Preprocessor1_M...

	<int>		<dbl>	<chr>
1	3		0.0316	Preprocessor1_Model6

Step 9: Finalize the model

```
# Final workflow
xgb_final_workflow <- xgb_workflow %>%
  finalize_workflow(best_xgb)
xgb_final_workflow
```

== Workflow ==

Preprocessor: Recipe

Model: boost\_tree()

— Preprocessor —

5 Recipe Steps

- step\_impute\_median()
- step\_impute\_mode()
- step\_dummy()
- step\_zv()
- step\_normalize()

— Model —

Boosted Tree Model Specification (classification)

Main Arguments:

```
trees = 1000
tree_depth = 3
learn_rate = 0.0316227766016838
```

Computational engine: xgboost

```
if (file.exists("xgb_final_fit.rds")) {
  xgb_final_fit <- read_rds("xgb_final_fit.rds")
  xgb_final_fit
} else {

  # fit the final model on the whole training set
  xgb_final_fit <- xgb_final_workflow %>%
    last_fit(data_split)

  xgb_final_fit %>%
    write_rds("xgb_final_fit.rds")

  xgb_final_fit
}
```

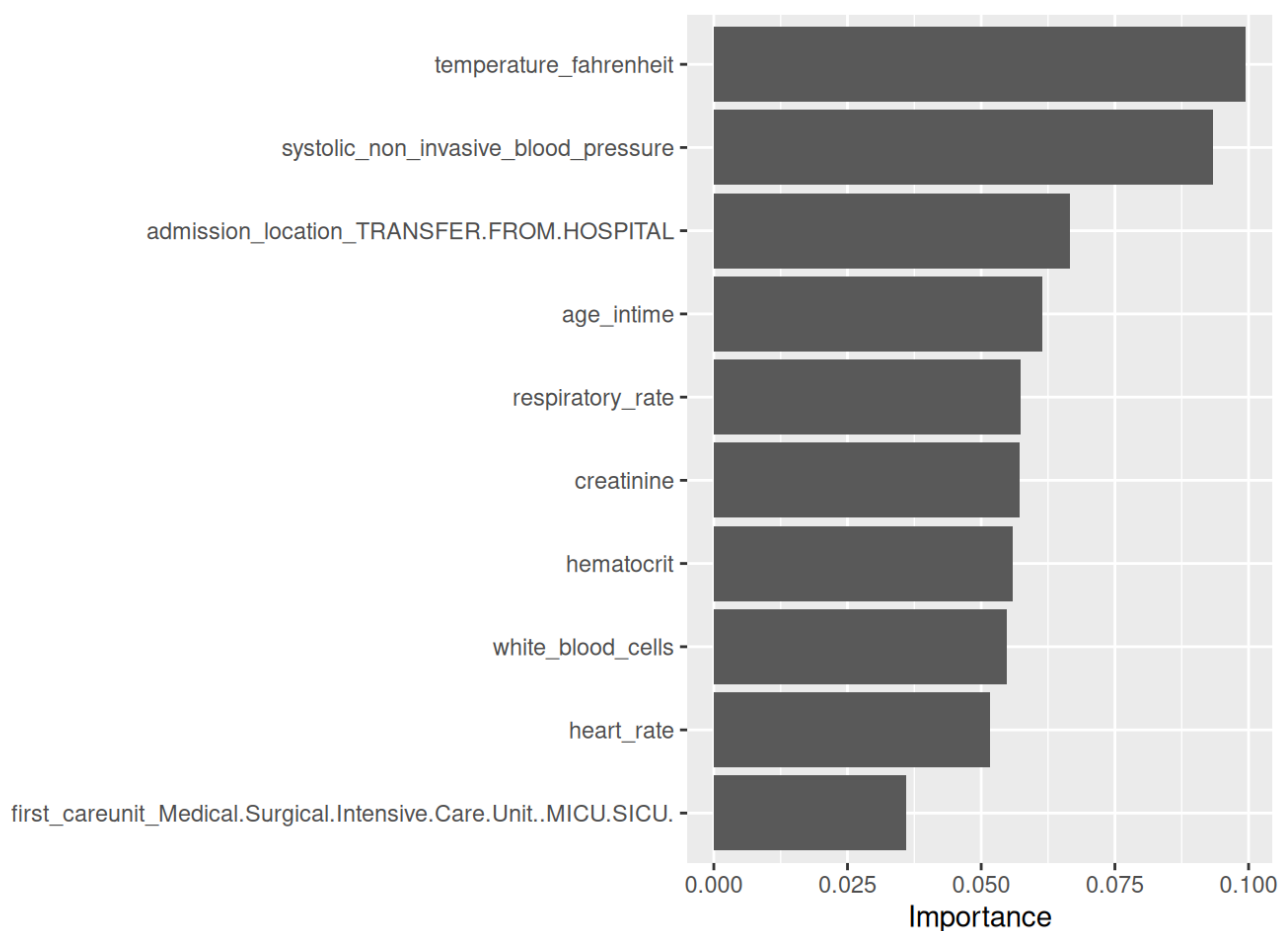
```
# Resampling results
# Manual resampling
# A tibble: 1 × 6
  splits          id      .metrics .notes  .predictions .workflow
  <list>         <chr>    <list>  <list>  <list>       <list>
1 <split [47221/47223]> train/test sp... <tibble> <tibble> <tibble>    <workflow>
```

```
# Test metrics
xgb_final_fit |>
  collect_metrics()
```

```
# A tibble: 3 × 4
  .metric      .estimator .estimate .config
  <chr>       <chr>         <dbl> <chr>
1 accuracy    binary         0.604 Preprocessor1_Model11
2 roc_auc     binary         0.647 Preprocessor1_Model11
3 brier_class binary         0.233 Preprocessor1_Model11
```

Step 10: Plot the variable importance

```
xgb_final_fit %>%
  extract_fit_engine() %>%
  vip::vip() %>%
  print()
```



Summary of the XGBoosting model: The model has accuracy of 0.6036677, meaning that the model can predict 60.4% of the test dataset correctly. The AUC of the model is 0.6465419, which means that the final XGBoosting model can correctly classify 64.7% of the test dataset. At this point, we can see that both the random forest model and XGBoosting model have better performance in both accuracy and AUC than the logistic model. The performance of random forest model and XGBoosting model are close to each other, with random forest has slightly higher accuracy and AUC.

The 10 most important predictors are listed in the variable importance plot. The most important variable in XGBoosting is temperature fahrenheit.

## Model Stacking

Step 1: Set up the cross-validation folds to be shared by 3 models using in the stack model

```
set.seed(203)
folds <- vfold_cv(train_data, v = 3)
```

Step 2: Base models. Here we use logistic regression with elastic net regularization, Random Forest, and XGBoosting. To shorten the running time, I changed the grid of each model to smaller numbers to ensure the system won't crash.

```
# Logistic regression with elastic net regularization
logit_mod <-
  logistic_reg(
    penalty = tune(),
    mixture = tune()
  ) |>
  set_engine("glmnet", standardize = TRUE)
logit_mod
```

### Logistic Regression Model Specification (classification)

Main Arguments:

```
penalty = tune()
mixture = tune()
```

Engine-Specific Arguments:

```
standardize = TRUE
```

Computational engine: glmnet

```
logit_wf <- workflow() |>
  add_recipe(logit_recipe) |>
  add_model(logit_mod)
logit_wf
```

== Workflow ==  
Preprocessor: Recipe

Model: `logistic_reg()`

— Preprocessor —

5 Recipe Steps

- `step_impute_median()`
- `step_impute_mode()`
- `step_dummy()`
- `step_zv()`
- `step_normalize()`

— Model —

Logistic Regression Model Specification (classification)

Main Arguments:

`penalty = tune()`  
`mixture = tune()`

Engine-Specific Arguments:

`standardize = TRUE`

Computational engine: `glmnet`

```
logit_stack_grid <- grid_regular(  
  penalty(range = c(-6, 3)),  
  mixture(),  
  levels = c(5, 5)  
)  
  
suppressMessages(suppressWarnings({  
  if (file.exists("logit_stack.rds")) {  
    logit_stack <- read_rds("logit_stack.rds")  
    logit_stack  
  
  } else {  
    (logit_stack <- logit_wf |>  
      tune_grid(  
        resamples = folds,  
        grid = logit_stack_grid,  
        metrics = metric_set(roc_auc, accuracy),  
        control = control_stack_grid()  
      )) |>  
    system.time()  
  
    logit_stack |>  
      write_rds("logit_stack.rds")  
  
    logit_stack  
  }  
}))
```

```
# Tuning results
# 3-fold cross-validation
# A tibble: 3 × 5
  splits          id    .metrics          .notes          .predictions
  <list>         <chr> <list>          <list>          <list>
1 <split [31480/15741]> Fold1 <tibble [50 × 6]> <tibble [0 × 3]> <tibble>
2 <split [31481/15740]> Fold2 <tibble [50 × 6]> <tibble [0 × 3]> <tibble>
3 <split [31481/15740]> Fold3 <tibble [50 × 6]> <tibble [0 × 3]> <tibble>
```

```
# Random forest
rf_mod <-
  rand_forest(
    mode = "classification",
    mtry = tune(),
    trees = tune()
  ) |>
  set_engine("ranger")
rf_mod
```

Random Forest Model Specification (classification)

Main Arguments:

```
mtry = tune()
trees = tune()
```

Computational engine: ranger

```
rf_wf <- workflow() |>
  add_recipe(rf_recipe) |>
  add_model(rf_mod)
rf_wf
```

== Workflow ==

Preprocessor: Recipe

Model: rand\_forest()

— Preprocessor —

5 Recipe Steps

- step\_impute\_median()
- step\_impute\_mode()
- step\_dummy()
- step\_zv()
- step\_normalize()

— Model —

Random Forest Model Specification (classification)

Main Arguments:

```
mtry = tune()
trees = tune()
```

Computational engine: ranger

```
rf_stack_grid <- grid_regular(
  trees(range = c(200L, 500L)),
  mtry(range = c(1L, 5L)),
  levels = c(5, 2)
)

if (file.exists("rf_stack.rds")) {
  rf_stack <- read_rds("rf_stack.rds")
  rf_stack
} else {
  (rf_stack <- rf_wf %>%
    tune_grid(
      resamples = folds,
      grid = rf_stack_grid,
      metrics = metric_set(roc_auc, accuracy),
      control = control_stack_grid()
    )) %>%
  system.time()

  rf_stack %>%
    write_rds("rf_stack.rds")

  rf_stack
}
```

# Tuning results

# 3-fold cross-validation

# A tibble: 3 × 5

	splits	id	.metrics	.notes	.predictions
	<list>	<chr>	<list>	<list>	<list>
1	<split [31480/15741]>	Fold1	<tibble [20 × 6]>	<tibble [0 × 3]>	<tibble>
2	<split [31481/15740]>	Fold2	<tibble [20 × 6]>	<tibble [0 × 3]>	<tibble>
3	<split [31481/15740]>	Fold3	<tibble [20 × 6]>	<tibble [0 × 3]>	<tibble>

```
# XGBoosting
gb_mod <-
  boost_tree(
    mode = "classification",
    trees = 1000,
    tree_depth = tune(),
    learn_rate = tune()
  ) |>
```

```
set_engine("xgboost")
gb_mod
```

### Boosted Tree Model Specification (classification)

#### Main Arguments:

```
trees = 1000
tree_depth = tune()
learn_rate = tune()
```

Computational engine: xgboost

```
gb_wf <- workflow() |>
  add_recipe(xgb_recipe) |>
  add_model(gb_mod)
gb_wf
```

#### == Workflow ==

Preprocessor: Recipe

Model: boost\_tree()

#### — Preprocessor —

5 Recipe Steps

- step\_impute\_median()
- step\_impute\_mode()
- step\_dummy()
- step\_zv()
- step\_normalize()

#### — Model —

### Boosted Tree Model Specification (classification)

#### Main Arguments:

```
trees = 1000
tree_depth = tune()
learn_rate = tune()
```

Computational engine: xgboost

```
gb_stack_grid <- grid_regular(
  tree_depth(range = c(1L, 3L)),
  learn_rate(range = c(-3, 1), trans = log10_trans()),
  levels = c(3, 3)
)
gb_stack_grid
```

# A tibble: 9 × 2

```
tree_depth learn_rate
```



	<int>	<dbl>
1	1	0.001
2	2	0.001
3	3	0.001
4	1	0.1
5	2	0.1
6	3	0.1
7	1	10
8	2	10
9	3	10

```

if (file.exists("gb_stack.rds")) {
  gb_stack <- read_rds("gb_stack.rds")
  gb_stack

} else {
  (gb_stack <- gb_wf %>%
    tune_grid(
      resamples = folds,
      grid = gb_stack_grid,
      metrics = metric_set(roc_auc, accuracy),
      control = control_stack_grid()
    )) %>%
    system.time()

  gb_stack %>%
    write_rds("gb_stack.rds")

  gb_stack
}

```

# Tuning results

# 3-fold cross-validation

# A tibble: 3 × 5

	splits	id	.metrics	.notes	.predictions
	<list>	<chr>	<list>	<list>	<list>
1	<split [31480/15741]>	Fold1	<tibble [18 × 6]>	<tibble [0 × 3]>	<tibble>
2	<split [31481/15740]>	Fold2	<tibble [18 × 6]>	<tibble [0 × 3]>	<tibble>
3	<split [31481/15740]>	Fold3	<tibble [18 × 6]>	<tibble [0 × 3]>	<tibble>

Step 3: Build the stacked ensemble

```

if (file.exists("stacks.rds")) {
  model_st <- read_rds("stacks.rds")
  model_st
} else {
  suppressWarnings({
    model_st <-
      stacks() |>

    # add candidate models

```

```

add_candidates(logit_stack) |>
add_candidates(rf_stack) |>
add_candidates(gb_stack) |>

# determine how to combine their predictions
blend_predictions(
  penalty = 10^(-5:2),
  metrics = c("roc_auc"),

  # set the number of resamples to 3 to reduce computation time
  times = 3) |>

# fit the candidates with nonzero stacking coefficients
fit_members()

model_st |> write_rds("stacks.rds")

model_st
})
}

```

— A stacked ensemble model —————

Out of 31 possible candidate members, the ensemble retained 9.

Penalty: 0.001.

Mixture: 1.

The 9 highest weighted member classes are:


```

# A tibble: 9 × 3
  member                type      weight
  <chr>                <chr>    <dbl>
1 .pred_TRUE_rf_stack_1_09 rand_forest 1.26
2 .pred_TRUE_rf_stack_1_06 rand_forest 1.01
3 .pred_TRUE_gb_stack_1_6  boost_tree 0.984
4 .pred_TRUE_rf_stack_1_07 rand_forest 0.816
5 .pred_TRUE_gb_stack_1_5  boost_tree 0.702
6 .pred_TRUE_rf_stack_1_10 rand_forest 0.585
7 .pred_TRUE_rf_stack_1_08 rand_forest 0.566
8 .pred_TRUE_gb_stack_1_8  boost_tree 0.0722
9 .pred_TRUE_gb_stack_1_7  boost_tree 0.0229


```

Step 4: Plot the stacked results

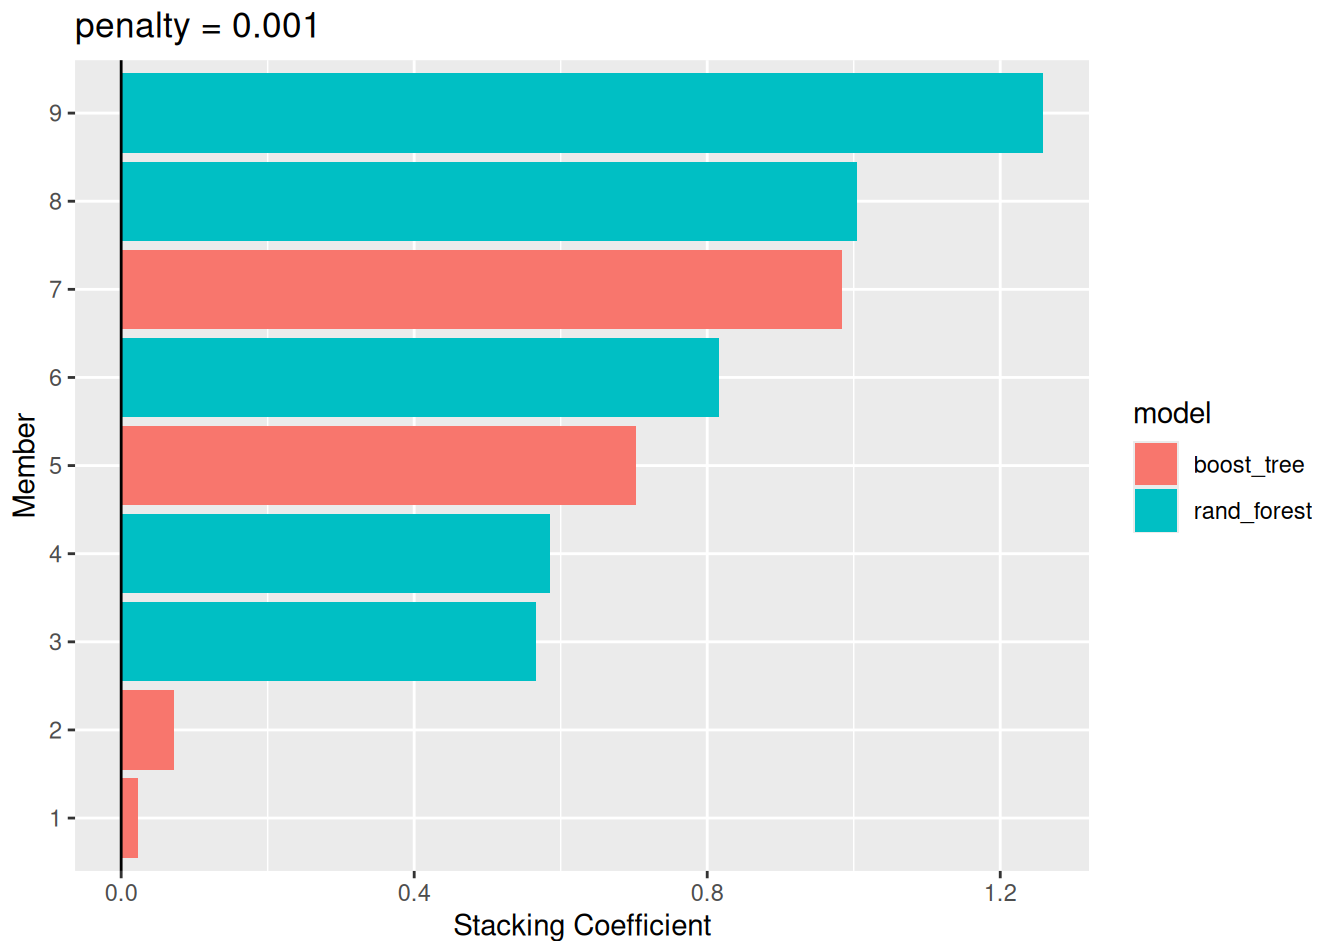
```
autoplot(model_st)
```



```
# To show the relationship more directly  
autoplot(model_st, type = "members")
```



```
# To see the top models  
autoplot(model_st, type = "weights")
```



Step 5: To identify which model configurations were assigned what stacking coefficients

```
# Coefficients of random forest model
collect_parameters(model_st, "rf_stack")
```

# A tibble: 10 × 5

member	mtry	trees	terms	coef
<chr>	<int>	<int>	<chr>	<dbl>
1 rf_stack_1_01	1	200	.pred_TRUE_rf_stack_1_01	0
2 rf_stack_1_02	1	275	.pred_TRUE_rf_stack_1_02	0
3 rf_stack_1_03	1	350	.pred_TRUE_rf_stack_1_03	0
4 rf_stack_1_04	1	425	.pred_TRUE_rf_stack_1_04	0
5 rf_stack_1_05	1	500	.pred_TRUE_rf_stack_1_05	0
6 rf_stack_1_06	5	200	.pred_TRUE_rf_stack_1_06	1.01
7 rf_stack_1_07	5	275	.pred_TRUE_rf_stack_1_07	0.816
8 rf_stack_1_08	5	350	.pred_TRUE_rf_stack_1_08	0.566
9 rf_stack_1_09	5	425	.pred_TRUE_rf_stack_1_09	1.26
10 rf_stack_1_10	5	500	.pred_TRUE_rf_stack_1_10	0.585

```
# Coefficients of XGBoosting model
collect_parameters(model_st, "gb_stack")
```

```
# A tibble: 9 × 5
```

member	tree_depth	learn_rate	terms	coef
<chr>	<int>	<dbl>	<chr>	<dbl>
1 gb_stack_1_1	1	0.001	.pred_TRUE_gb_stack_1_1	0
2 gb_stack_1_2	2	0.001	.pred_TRUE_gb_stack_1_2	0
3 gb_stack_1_3	3	0.001	.pred_TRUE_gb_stack_1_3	0
4 gb_stack_1_4	1	0.1	.pred_TRUE_gb_stack_1_4	0
5 gb_stack_1_5	2	0.1	.pred_TRUE_gb_stack_1_5	0.702
6 gb_stack_1_6	3	0.1	.pred_TRUE_gb_stack_1_6	0.984
7 gb_stack_1_7	1	10	.pred_TRUE_gb_stack_1_7	0.0229
8 gb_stack_1_8	2	10	.pred_TRUE_gb_stack_1_8	0.0722
9 gb_stack_1_9	3	10	.pred_TRUE_gb_stack_1_9	0

```
# Coefficients of Logistic regression model
collect_parameters(model_st, "logit_stack")
```

```
# A tibble: 12 × 5
```

member	penalty	mixture	terms	coef
<chr>	<dbl>	<dbl>	<chr>	<dbl>
1 logit_stack_1_01	0.000001	0	.pred_TRUE_logit_stack_1_01	0
2 logit_stack_1_03	0.0316	0	.pred_TRUE_logit_stack_1_03	0
3 logit_stack_1_04	5.62	0	.pred_TRUE_logit_stack_1_04	0
4 logit_stack_1_06	0.000001	0.25	.pred_TRUE_logit_stack_1_06	0
5 logit_stack_1_08	0.0316	0.25	.pred_TRUE_logit_stack_1_08	0
6 logit_stack_1_09	5.62	0.25	.pred_TRUE_logit_stack_1_09	0
7 logit_stack_1_11	0.000001	0.5	.pred_TRUE_logit_stack_1_11	0
8 logit_stack_1_13	0.0316	0.5	.pred_TRUE_logit_stack_1_13	0
9 logit_stack_1_16	0.000001	0.75	.pred_TRUE_logit_stack_1_16	0
10 logit_stack_1_18	0.0316	0.75	.pred_TRUE_logit_stack_1_18	0
11 logit_stack_1_21	0.000001	1	.pred_TRUE_logit_stack_1_21	0
12 logit_stack_1_23	0.0316	1	.pred_TRUE_logit_stack_1_23	0

Step 6: Finalization Apply the model on the test data and output the final classification

```
if (file.exists("final_classification.rds")) {
  final_classification <- read_rds("final_classification.rds")

  final_classification

} else {
  final_classification <- test_data |>
    bind_cols(predict(model_st, test_data, type = "prob")) |>
    print(width = Inf)

  final_classification

  final_classification |>
    write_rds("final_classification.rds")
}
```

```
# A tibble: 47,223 × 28
  first_careunit      admission_type admission_location insurance language
  <fct>              <fct>          <fct>          <fct>    <fct>
1 Medical Intensive Care ... EW EMER.      EMERGENCY ROOM    Medicaid English
2 Medical/Surgical Intens... EW EMER.      Other              Private   English
3 Cardiac Vascular Intens... SURGICAL SAME... PHYSICIAN REFERRAL Medicare English
4 Coronary Care Unit (CCU) OBSERVATION A... PHYSICIAN REFERRAL Medicaid English
5 Medical Intensive Care ... EW EMER.      EMERGENCY ROOM    Medicare English
6 Medical/Surgical Intens... EW EMER.      EMERGENCY ROOM    Medicare English
7 Medical Intensive Care ... EW EMER.      EMERGENCY ROOM    Medicare English
8 Coronary Care Unit (CCU) URGENT              TRANSFER FROM HOS... Medicare English
9 Coronary Care Unit (CCU) EW EMER.      TRANSFER FROM HOS... Private   English
10 Cardiac Vascular Intens... OBSERVATION A... PHYSICIAN REFERRAL Private   English
# i 47,213 more rows
# i 23 more variables: marital_status <fct>, race <fct>, gender <chr>,
#   anchor_age <int>, anchor_year <int>, anchor_year_group <chr>,
#   bicarbonate <dbl>, chloride <dbl>, creatinine <dbl>, glucose <dbl>,
#   potassium <dbl>, sodium <dbl>, hematocrit <dbl>, white_blood_cells <dbl>,
#   heart_rate <dbl>, systolic_non_invasive_blood_pressure <dbl>,
#   diastolic_non_invasive_blood_pressure <dbl>, ...
```

Compute the ROC AUC and accuracy of the final classification

```
# ROC AUC
yardstick::roc_auc(
  final_classification,
  truth = los_long,
  contains(".pred_FALSE")
)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>   <chr>         <dbl>
1 roc_auc binary      0.653
```

```
# Accuracy
final_classification <- final_classification %>%
  mutate(.pred_class = as.factor(
    ifelse(.pred_TRUE > .pred_FALSE, "TRUE", "FALSE")))

yardstick::accuracy(
  final_classification,
  truth = los_long,
  estimate = .pred_class
)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>   <chr>         <dbl>
1 accuracy binary      0.609
```

Use the members argument to generate predictions from each of the ensemble members

```
if (file.exists("mimic_pred.rds")) {
  mimic_pred <- read_rds("mimic_pred.rds")
} else {
  mimic_pred <-
    test_data |>
    select(los_long) |>
    bind_cols(
      predict(
        model_st,
        test_data,
        type = "class",
        members = TRUE
      )
    ) |>
    print(width = Inf)

  write_rds(mimic_pred, "mimic_pred.rds")
}
```

mimic\_pred

```
# A tibble: 47,223 × 11
  los_long .pred_class .pred_class_rf_stack_1_06 .pred_class_rf_stack_1_07
  <fct>    <fct>        <fct>          <fct>
1 FALSE   TRUE          TRUE           TRUE
2 FALSE   FALSE        FALSE          FALSE
3 FALSE   FALSE        FALSE          FALSE
4 TRUE    TRUE          TRUE           TRUE
5 FALSE   TRUE          TRUE           TRUE
6 TRUE    TRUE          TRUE           TRUE
7 TRUE    FALSE        FALSE          FALSE
8 TRUE    TRUE          TRUE           TRUE
9 TRUE    FALSE        FALSE          TRUE
10 FALSE  TRUE          TRUE           TRUE
# i 47,213 more rows
# i 7 more variables: .pred_class_rf_stack_1_08 <fct>,
#   .pred_class_rf_stack_1_09 <fct>, .pred_class_rf_stack_1_10 <fct>,
#   .pred_class_gb_stack_1_7 <fct>, .pred_class_gb_stack_1_5 <fct>,
#   .pred_class_gb_stack_1_8 <fct>, .pred_class_gb_stack_1_6 <fct>
```

```
# Get the mean of the predicted classes for each model
map(
  colnames(mimic_pred),
  ~mean(mimic_pred$los_long == pull(mimic_pred, .x))
) |>
set_names(colnames(mimic_pred)) |>
```



```
as_tibble() |>
pivot_longer(c(everything(), -los_long))
```

```
# A tibble: 10 × 3
  los_long name          value
  <dbl> <chr>          <dbl>
1      1 .pred_class      0.609
2      1 .pred_class_rf_stack_1_06 0.603
3      1 .pred_class_rf_stack_1_07 0.604
4      1 .pred_class_rf_stack_1_08 0.606
5      1 .pred_class_rf_stack_1_09 0.604
6      1 .pred_class_rf_stack_1_10 0.605
7      1 .pred_class_gb_stack_1_7  0.454
8      1 .pred_class_gb_stack_1_5  0.603
9      1 .pred_class_gb_stack_1_8  0.448
10     1 .pred_class_gb_stack_1_6  0.605
```

#### 4. Compare model classification performance on the test set. Report both the area under ROC curve and accuracy for each machine learning algorithm and the model stacking. Interpret the results. What are the most important features in predicting long ICU stays? How do the models compare in terms of performance and interpretability?

Report the information of accuracy and AUC of each machine learning algorithm and the model stacking

```
# Logistic regression with elastic net regularization
logit_final_fit %>%
  collect_metrics()
```

```
# A tibble: 3 × 4
  .metric .estimator .estimate .config
  <chr>    <chr>        <dbl> <chr>
1 accuracy binary        0.578 Preprocessor1_Model1
2 roc_auc  binary        0.608 Preprocessor1_Model1
3 brier_class binary        0.241 Preprocessor1_Model1
```

```
# Random forest
rf_final_fit %>%
  collect_metrics()
```

```
# A tibble: 3 × 4
  .metric .estimator .estimate .config
  <chr>    <chr>        <dbl> <chr>
1 accuracy binary        0.606 Preprocessor1_Model1
2 roc_auc  binary        0.647 Preprocessor1_Model1
3 brier_class binary        0.234 Preprocessor1_Model1
```

```
# XGBoosting
xgb_final_fit %>%
  collect_metrics()
```

```
# A tibble: 3 × 4
  .metric      .estimator .estimate .config
  <chr>        <chr>      <dbl> <chr>
1 accuracy    binary      0.604 Preprocessor1_Model1
2 roc_auc     binary      0.647 Preprocessor1_Model1
3 brier_class binary      0.233 Preprocessor1_Model1
```

```
# Model stacking
## Accuracy
final_classification <- final_classification %>%
  mutate(.pred_class = as.factor(
    ifelse(.pred_TRUE > .pred_FALSE, "TRUE", "FALSE")))

yardstick::accuracy(
  final_classification,
  truth = los_long,
  estimate = .pred_class
)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>    <chr>      <dbl>
1 accuracy binary      0.609
```

```
## ROC AUC
yardstick::roc_auc(
  final_classification,
  truth = los_long,
  contains(".pred_FALSE")
)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>    <chr>      <dbl>
1 roc_auc binary      0.653
```

It can be seen that the accuracy and AUC of the logistic regression model are 0.5782775 and 0.6081425 respectively, the accuracy and AUC of the random forest model are 0.6060818 and 0.6474515 respectively, and the accuracy and AUC of the XGBoosting model are 0.6036677 and 0.6465419 respectively. Based on the above results, we can see that for individual models, random forest has the highest accuracy and AUC and thus it is the best-performing model among the three machine learning algorithms, and it has the greatest weight in the stacked ensemble, followed by the XGBoosting model. In order to speed up the processing and prevent the system from crashing, I selected a smaller fold of cross-validation (3) than the ones for single models (5), and I also reduced the grids. Because of this reason, there are some duplicated candidates that are automatically removed by the system,

and the stacked results shows no information on logistic model. But this should not be a problem since the logistic model is the least efficient one among the three models.

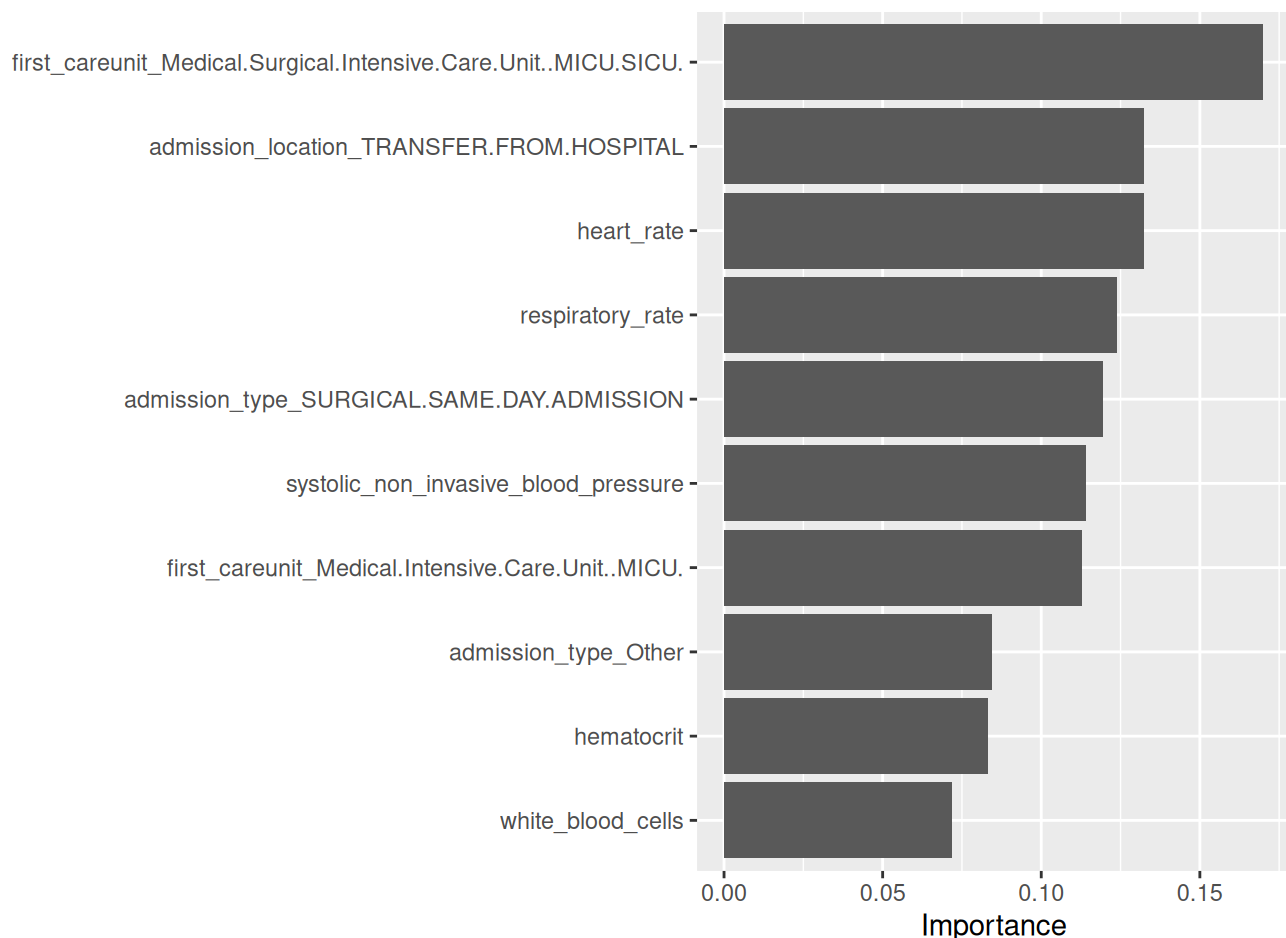
The accuracy and AUC of the final classification of the stacked model are 0.6091947 and 0.6530818 respectively, meaning that the probability of the model ranking a randomly chosen positive observation higher than a randomly chosen negative observation is 0.6530818 and the probability of correct prediction is 0.6091947.

Turning to the coefficients of each model within the stacked ensemble, the random forest model occupied the top 5 values and the XGBoosting model takes the remaining 4 positions. No information about the logistic model is shown because of the reason I mentioned earlier.


To summarize, the random forest model has the best performance in the prediction, followed by the XGBoosting model, and logistic model being the least efficient one.

## Report the most important features in predicting long ICU stays

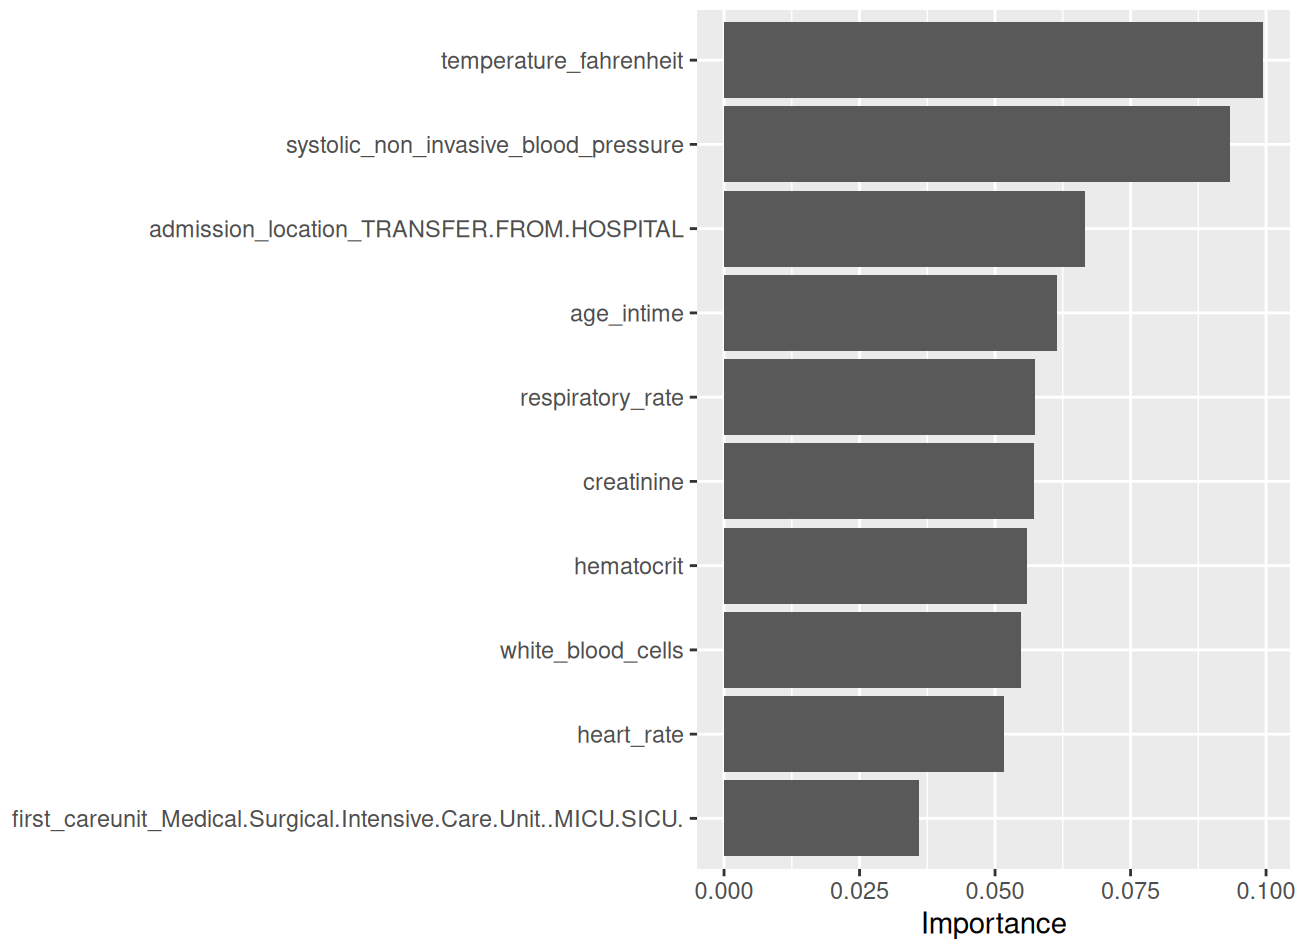
```
# Logistic regression with elastic net regularization
logit_final_fit %>%
  extract_fit_parsnip() %>%
  vip::vip() %>%
  print()
```



```
# Random forest
rf_final_fit %>%
  extract_fit_engine() %>%
  vip::vip() %>%
  print()
```



```
# XGBoosting
xgb_final_fit %>%
  extract_fit_engine() %>%
  vip::vip() %>%
  print()
```



As shown in the above plots, the most important feature in predicting long ICU stays for logistic regression model is first careunit. For random forest model, the most important feature is systolic noninvasive blood pressure. If we look close enough, we will find that heart rate also plays a very important role in prediction. In fact, for this model, there are 7 variables (systolic noninvasive blood pressure, heart rate, white blood cells, hematocrit, temperature fahrenheit, diastolic noninvasive blood pressure, and respiratory rate) exit the importance of 750. For XGBoosting model, the most important variable is temperature fahrenheit. ##### Summary of performance and interpretability In general, random forest model has the best performance, followed by XGBoosting model, and logistic regression model the least efficient. But in terms of intepretability, logistic regression model is the best.